
Course: Laboratory Practice III

Course Code: 410246

Group Members :

Pravin Jain

Onkar Kulkarni

Class: BE

Mini-Project Title: Build a machine learning model that predicts the type of people who survived the Titanic shipwreck using passenger data (i.e. name, age, gender, socio-economic class, etc.).

Dataset Link: <https://www.kaggle.com/competitions/titanic/data>

```
import numpy as np

# data processing
import pandas as pd

# data visualization
import seaborn as sns
%matplotlib inline
from matplotlib import pyplot as plt
from matplotlib import style

# Algorithms
from sklearn import linear_model
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import Perceptron
from sklearn.linear_model import SGDClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC, LinearSVC
from sklearn.naive_bayes import GaussianNB

test_df = pd.read_csv("test.csv")
train_df = pd.read_csv("train.csv")

train_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId      891 non-null    int64
1   Survived         891 non-null    int64
```

```

2  Pclass      891 non-null    int64
3  Name        891 non-null    object
4  Sex          891 non-null    object
5  Age          714 non-null    float64
6  SibSp        891 non-null    int64
7  Parch        891 non-null    int64
8  Ticket       891 non-null    object
9  Fare         891 non-null    float64
10 Cabin        204 non-null    object
11 Embarked     889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB

```


```
train_df.head(8)
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.250
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.283
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.925
3	4	1	1	Fraenkel, Leo. A.	male	35.0	1	0	3101538	53.100
4	5	0	3	Allen, Mr. William Henry	male	29.0	0	0	241601	8.050
5	6	0	1	Meyer, Mr. Jacob	male	65.0	0	0	9140	51.000
6	7	0	3	Wright, Mrs. Mary	female	31.0	0	0	2101282	7.250

```

total = train_df.isnull().sum().sort_values(ascending=False)
percent_1 = train_df.isnull().sum()/train_df.isnull().count()*100
percent_2 = (round(percent_1, 1)).sort_values(ascending=False)
missing_data = pd.concat([total, percent_2], axis=1, keys=['Total', '%'])
missing_data.head(5)

```


	Total	%	
Cabin	687	77.1	
Age	177	19.9	
Embarked	2	0.2	
PassengerId	0	0.0	
Survived	0	0.0	

```
train_df.columns.values
```

```
array(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',
      'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'], dtype=object)
```


```
survived = 'survived'
not_survived = 'not survived'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(10, 4))
women = train_df[train_df['Sex']=='female']
men = train_df[train_df['Sex']=='male']
ax = sns.distplot(women[women['Survived']==1].Age.dropna(), bins=18, label = survived, ax = axes[0])
ax = sns.distplot(women[women['Survived']==0].Age.dropna(), bins=40, label = not_survived, ax = axes[0])
ax.legend()
ax.set_title('Female')
ax = sns.distplot(men[men['Survived']==1].Age.dropna(), bins=18, label = survived, ax = axes[1])
ax = sns.distplot(men[men['Survived']==0].Age.dropna(), bins=40, label = not_survived, ax = axes[1])
ax.legend()
_ = ax.set_title('Male')
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: FutureWarning: `distplot` is a deprecated alias for `displot`. Please adjust your code to use `displot` instead.
warnings.warn(msg, FutureWarning)
```




```
FacetGrid = sns.FacetGrid(train_df, row='Embarked', size=4.5, aspect=1.6)
FacetGrid.map(sns.pointplot, 'Pclass', 'Survived', 'Sex', palette=None, order=None, hue_order=None)
FacetGrid.add_legend()
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/axisgrid.py:337: UserWarning: The `size`  
warnings.warn(msg, UserWarning)  
<seaborn.axisgrid.FacetGrid at 0x7f9c5b3e21d0>
```



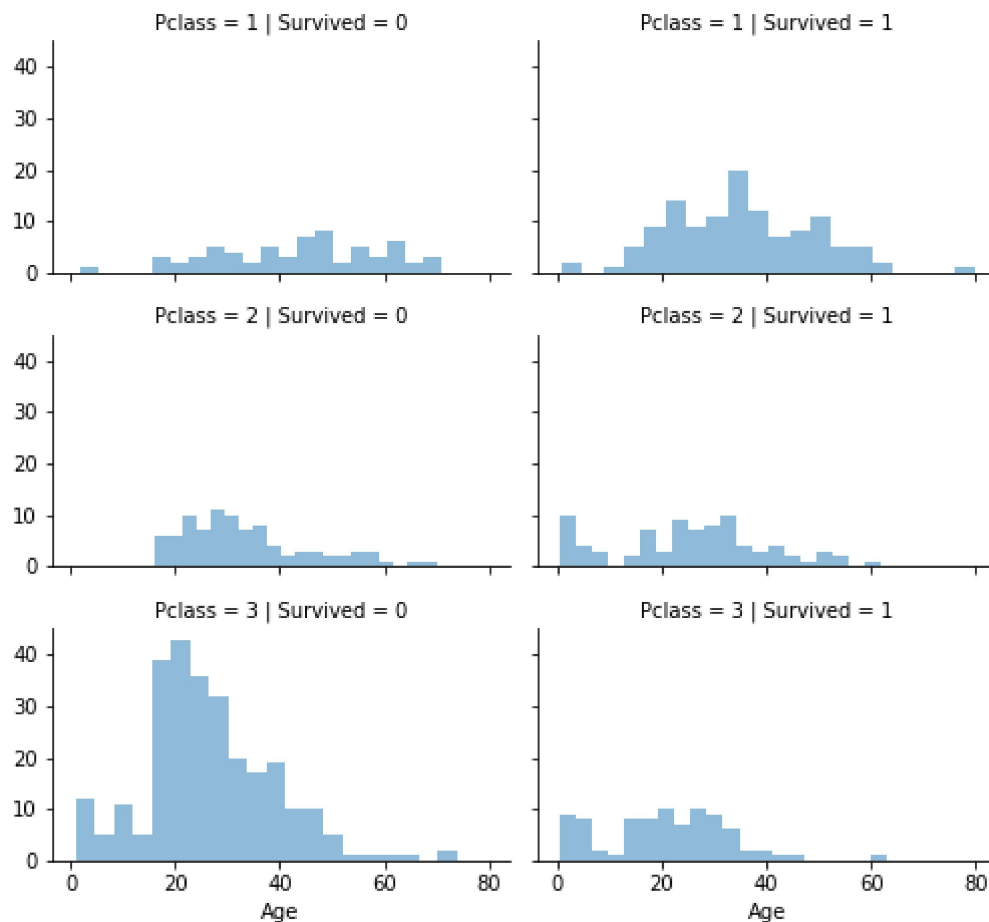
```
sns.barplot(x='Pclass', y='Survived', data=train_df)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f9c5b3db810>
```



```
grid = sns.FacetGrid(train_df, col='Survived', row='Pclass', size=2.2, aspect=1.6)
grid.map(plt.hist, 'Age', alpha=.5, bins=20)
grid.add_legend();
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/axisgrid.py:337: UserWarning: The `size`
warnings.warn(msg, UserWarning)
```



```
data = [train_df, test_df]
for dataset in data:
    dataset['relatives'] = dataset['SibSp'] + dataset['Parch']
    dataset.loc[dataset['relatives'] > 0, 'not_alone'] = 0
    dataset.loc[dataset['relatives'] == 0, 'not_alone'] = 1
    dataset['not_alone'] = dataset['not_alone'].astype(int)
train_df['not_alone'].value_counts()
```

```
1    537
0    354
Name: not_alone, dtype: int64
```

DATA PRE-PROCESSING

```
train_df = train_df.drop(['PassengerId'], axis=1)
```

```
import re
```

```
deck = {"A": 1, "B": 2, "C": 3, "D": 4, "E": 5, "F": 6, "G": 7, "U": 8}
```

```
data = [train_df, test_df]
```

```
for dataset in data:
```

```
    dataset['Cabin'] = dataset['Cabin'].fillna("U0")
```

```
    dataset['Deck'] = dataset['Cabin'].map(lambda x: re.compile("([a-zA-Z]+)").search(x).group(1))
```

```
    dataset['Deck'] = dataset['Deck'].map(deck)
```

```
    dataset['Deck'] = dataset['Deck'].fillna(0)
```

```
    dataset['Deck'] = dataset['Deck'].astype(int)
```

```
# we can now drop the cabin feature
```

```
train_df = train_df.drop(['Cabin'], axis=1)
```

```
test_df = test_df.drop(['Cabin'], axis=1)
```

```
data = [train_df, test_df]
```

```
for dataset in data:
```

```
    mean = train_df["Age"].mean()
```

```
    std = test_df["Age"].std()
```

```
    is_null = dataset["Age"].isnull().sum()
```

```
# compute random numbers between the mean, std and is_null
```

```
rand_age = np.random.randint(mean - std, mean + std, size = is_null)
```

```
# fill NaN values in Age column with random values generated
```

```
age_slice = dataset["Age"].copy()
```

```
age_slice[np.isnan(age_slice)] = rand_age
```

```
dataset["Age"] = age_slice
```

```
dataset["Age"] = train_df["Age"].astype(int)
```

```
train_df["Age"].isnull().sum()
```

```
0
```

```
train_df['Embarked'].describe()
```

```
count    889
```

```
unique     3
```

```
top         S
```

```
freq       644
```

```
Name: Embarked, dtype: object
```

```
common_value = 'S'
```

```
data = [train_df, test_df]
```

```
for dataset in data:
```

```
    dataset['Embarked'] = dataset['Embarked'].fillna(common_value)
```

CONVERTING FEATURES

```
train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 13 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Survived    891 non-null    int64
1   Pclass      891 non-null    int64
2   Name        891 non-null    object
3   Sex         891 non-null    object
4   Age         891 non-null    int64
5   SibSp       891 non-null    int64
6   Parch       891 non-null    int64
7   Ticket      891 non-null    object
8   Fare        891 non-null    float64
9   Embarked    891 non-null    object
10  relatives   891 non-null    int64
11  not_alone   891 non-null    int64
12  Deck        891 non-null    int64
dtypes: float64(1), int64(8), object(4)
memory usage: 90.6+ KB
```

```
data = [train_df, test_df]
```

```
for dataset in data:
```

```
    dataset['Fare'] = dataset['Fare'].fillna(0)
```

```
    dataset['Fare'] = dataset['Fare'].astype(int)
```

```
data = [train_df, test_df]
```

```
titles = {"Mr": 1, "Miss": 2, "Mrs": 3, "Master": 4, "Rare": 5}
```

```
for dataset in data:
```

```
    # extract titles
```

```
    dataset['Title'] = dataset.Name.str.extract('([A-Za-z]+)\.', expand=False)
```

```
    # replace titles with a more common title or as Rare
```

```
    dataset['Title'] = dataset['Title'].replace(['Lady', 'Countess', 'Capt', 'Col', 'Don', 'Dr',
                                                'Major', 'Rev', 'Sir', 'Jonkheer', 'Dona'], 'Rare')
```

```
    dataset['Title'] = dataset['Title'].replace('Mlle', 'Miss')
```

```
    dataset['Title'] = dataset['Title'].replace('Ms', 'Miss')
```

```
    dataset['Title'] = dataset['Title'].replace('Mme', 'Mrs')
```

```
    # convert titles into numbers
```

```

dataset['Title'] = dataset['Title'].map(titles)
# filling NaN with 0, to get safe
dataset['Title'] = dataset['Title'].fillna(0)
train_df = train_df.drop(['Name'], axis=1)
test_df = test_df.drop(['Name'], axis=1)

```

```

genders = {"male": 0, "female": 1}
data = [train_df, test_df]

```

```

for dataset in data:
    dataset['Sex'] = dataset['Sex'].map(genders)

```

```

train_df['Ticket'].describe()

```

```

count      891
unique      681
top        347082
freq         7
Name: Ticket, dtype: object

```

```

train_df = train_df.drop(['Ticket'], axis=1)
test_df = test_df.drop(['Ticket'], axis=1)

```

```

ports = {"S": 0, "C": 1, "Q": 2}
data = [train_df, test_df]

```

```

for dataset in data:
    dataset['Embarked'] = dataset['Embarked'].map(ports)

```

```

data = [train_df, test_df]

```

```

for dataset in data:
    dataset['Age'] = dataset['Age'].astype(int)
    dataset.loc[ dataset['Age'] <= 11, 'Age'] = 0
    dataset.loc[(dataset['Age'] > 11) & (dataset['Age'] <= 18), 'Age'] = 1
    dataset.loc[(dataset['Age'] > 18) & (dataset['Age'] <= 22), 'Age'] = 2
    dataset.loc[(dataset['Age'] > 22) & (dataset['Age'] <= 27), 'Age'] = 3
    dataset.loc[(dataset['Age'] > 27) & (dataset['Age'] <= 33), 'Age'] = 4
    dataset.loc[(dataset['Age'] > 33) & (dataset['Age'] <= 40), 'Age'] = 5
    dataset.loc[(dataset['Age'] > 40) & (dataset['Age'] <= 66), 'Age'] = 6
    dataset.loc[ dataset['Age'] > 66, 'Age'] = 6
train_df['Age'].value_counts()

```

```

4    164
6    158
5    153
3    138
2    114
1     96

```



```
0      68
Name: Age, dtype: int64
```

```
train_df.head(10)
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked	relatives	not_alone	Deck
0	0	3	0	2	1	0	7	0	1	0	8
1	1	1	1	5	1	0	71	1	1	0	3
2	1	3	1	3	0	0	7	0	0	1	8
3	1	1	1	5	1	0	53	0	1	0	3
4	0	3	0	5	0	0	8	0	0	1	8
5	0	3	0	1	0	0	8	2	0	1	8
6	0	1	0	6	0	0	51	0	0	1	5
7	0	3	0	0	3	1	21	0	4	0	8
8	1	3	1	3	0	2	11	0	2	0	8
9	1	2	1	1	1	0	30	1	1	0	8

```
data = [train_df, test_df]
```

```
for dataset in data:
```

```
    dataset.loc[ dataset['Fare'] <= 7.91, 'Fare'] = 0
    dataset.loc[(dataset['Fare'] > 7.91) & (dataset['Fare'] <= 14.454), 'Fare'] = 1
    dataset.loc[(dataset['Fare'] > 14.454) & (dataset['Fare'] <= 31), 'Fare'] = 2
    dataset.loc[(dataset['Fare'] > 31) & (dataset['Fare'] <= 99), 'Fare'] = 3
    dataset.loc[(dataset['Fare'] > 99) & (dataset['Fare'] <= 250), 'Fare'] = 4
    dataset.loc[ dataset['Fare'] > 250, 'Fare'] = 5
    dataset['Fare'] = dataset['Fare'].astype(int)
```

```
data = [train_df, test_df]
```

```
for dataset in data:
```

```
    dataset['Age_Class'] = dataset['Age'] * dataset['Pclass']
```

```
for dataset in data:
```

```
    dataset['Fare_Per_Person'] = dataset['Fare'] / (dataset['relatives'] + 1)
    dataset['Fare_Per_Person'] = dataset['Fare_Per_Person'].astype(int)
```

```
# Let's take a last look at the training set, before we start training the models.
```

```
train_df.head(10)
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked	relatives	not_alone	Deck
0	0	3	0	2	1	0	0	0	1	0	8
1	1	1	1	5	1	0	3	1	1	0	3
2	1	3	1	3	0	0	0	0	0	1	8
3	1	1	1	5	1	0	3	0	1	0	3
4	0	3	0	5	0	0	1	0	0	1	8
5	0	3	0	1	0	0	1	2	0	1	8
6	0	1	0	6	0	0	3	0	0	1	5
7	0	3	0	0	3	1	2	0	4	0	8

BUILDING MACHINE LEARNING MODELS

```

0      1      2      1      1      1      0      2      1      1      0      8

```

Random Forest

```

X_train = train_df.drop("Survived", axis=1)
Y_train = train_df["Survived"]
X_test  = test_df.drop("PassengerId", axis=1).copy()

random_forest = RandomForestClassifier(n_estimators=100)
random_forest.fit(X_train, Y_train)

Y_prediction = random_forest.predict(X_test)

random_forest.score(X_train, Y_train)
acc_random_forest = round(random_forest.score(X_train, Y_train) * 100, 2)

```

Logistic Regression

```

logreg = LogisticRegression()
logreg.fit(X_train, Y_train)

Y_pred = logreg.predict(X_test)

acc_log = round(logreg.score(X_train, Y_train) * 100, 2)

```

Naive Bayes

```

gaussian = GaussianNB()
gaussian.fit(X_train, Y_train)

```

```
Y_pred = gaussian.predict(X_test)
acc_gaussian = round(gaussian.score(X_train, Y_train) * 100, 2)
```

Linear Support Vector Machine

```
linear_svc = LinearSVC()
linear_svc.fit(X_train, Y_train)
```


```
Y_pred = linear_svc.predict(X_test)
```

```
acc_linear_svc = round(linear_svc.score(X_train, Y_train) * 100, 2)
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:1208: ConvergenceWarning: Li
ConvergenceWarning,
```



```
results = pd.DataFrame({
    'Model': ['Support Vector Machines', 'Logistic Regression',
             'Random Forest', 'Naive Bayes'],
    'Score': [acc_linear_svc, acc_log,
             acc_random_forest, acc_gaussian]})
result_df = results.sort_values(by='Score', ascending=False)
result_df = result_df.set_index('Score')
result_df.head(9)
```

	Model 
Score	
93.04	Random Forest
81.37	Logistic Regression
81.26	Support Vector Machines
78.68	Naive Bayes

```
from sklearn.model_selection import cross_val_score
rf = RandomForestClassifier(n_estimators=100)
scores = cross_val_score(rf, X_train, Y_train, cv=10, scoring = "accuracy")
print("Scores:", scores)
print("Mean:", scores.mean())
print("Standard Deviation:", scores.std())
```

```
Scores: [0.75555556 0.82022472 0.75280899 0.85393258 0.86516854 0.86516854
0.84269663 0.7752809 0.87640449 0.84269663]
Mean: 0.8249937578027465
Standard Deviation: 0.044566518656266645
```

Feature Importance

```
importances = pd.DataFrame({'feature':X_train.columns,'importance':np.round(random_forest.fea
importances = importances.sort_values('importance',ascending=False).set_index('feature')
importances.head(15)
```

	importance
feature	
Title	0.205
Sex	0.165
Age_Class	0.093
Deck	0.085
Age	0.075
Pclass	0.074
Fare	0.065
relatives	0.061
Embarked	0.057
Fare_Per_Person	0.044
SibSp	0.039
Parch	0.025
not_alone	0.012

```
importances.plot.bar()
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f9c580b2210>



```
train_df = train_df.drop("not_alone", axis=1)
test_df = test_df.drop("not_alone", axis=1)
```

```
train_df = train_df.drop("Parch", axis=1)
test_df = test_df.drop("Parch", axis=1)
```



```
random_forest = RandomForestClassifier(n_estimators=100, oob_score = True)
random_forest.fit(X_train, Y_train)
Y_prediction = random_forest.predict(X_test)
```

```
random_forest.score(X_train, Y_train)
```

```
acc_random_forest = round(random_forest.score(X_train, Y_train) * 100, 2)
print(round(acc_random_forest,2), "%")
```

93.04 %

Random Forest

```
random_forest = RandomForestClassifier(criterion = "gini",
                                       min_samples_leaf = 1,
                                       min_samples_split = 10,
                                       n_estimators=100,
                                       max_features='auto',
                                       oob_score=True,
                                       random_state=1,
                                       n_jobs=-1)
```

```
random_forest.fit(X_train, Y_train)
Y_prediction = random_forest.predict(X_test)
```

```
random_forest.score(X_train, Y_train)
```

```
print("oob score:", round(random_forest.oob_score_, 4)*100, "%")
```

oob score: 82.49 %

Evaluation

```
from sklearn.model_selection import cross_val_predict
from sklearn.metrics import confusion_matrix
predictions = cross_val_predict(random_forest, X_train, Y_train, cv=3)
confusion_matrix(Y_train, predictions)
```

array([[491, 58],

```
[ 97, 245]])
```

```
from sklearn.metrics import precision_score, recall_score
```

```
print("Precision:", precision_score(Y_train, predictions))
```

```
print("Recall:", recall_score(Y_train, predictions))
```

```
Precision: 0.8085808580858086
```

```
Recall: 0.716374269005848
```

```
from sklearn.metrics import f1_score
```

```
f1_score(Y_train, predictions)
```

```
0.7596899224806202
```

```
from sklearn.metrics import precision_recall_curve
```

```
# getting the probabilities of our predictions
```

```
y_scores = random_forest.predict_proba(X_train)
```

```
y_scores = y_scores[:,1]
```

```
precision, recall, threshold = precision_recall_curve(Y_train, y_scores)
```

```
def plot_precision_and_recall(precision, recall, threshold):
```

```
    plt.plot(threshold, precision[:-1], "r-", label="precision", linewidth=5)
```

```
    plt.plot(threshold, recall[:-1], "b", label="recall", linewidth=5)
```

```
    plt.xlabel("threshold", fontsize=19)
```

```
    plt.legend(loc="upper right", fontsize=19)
```

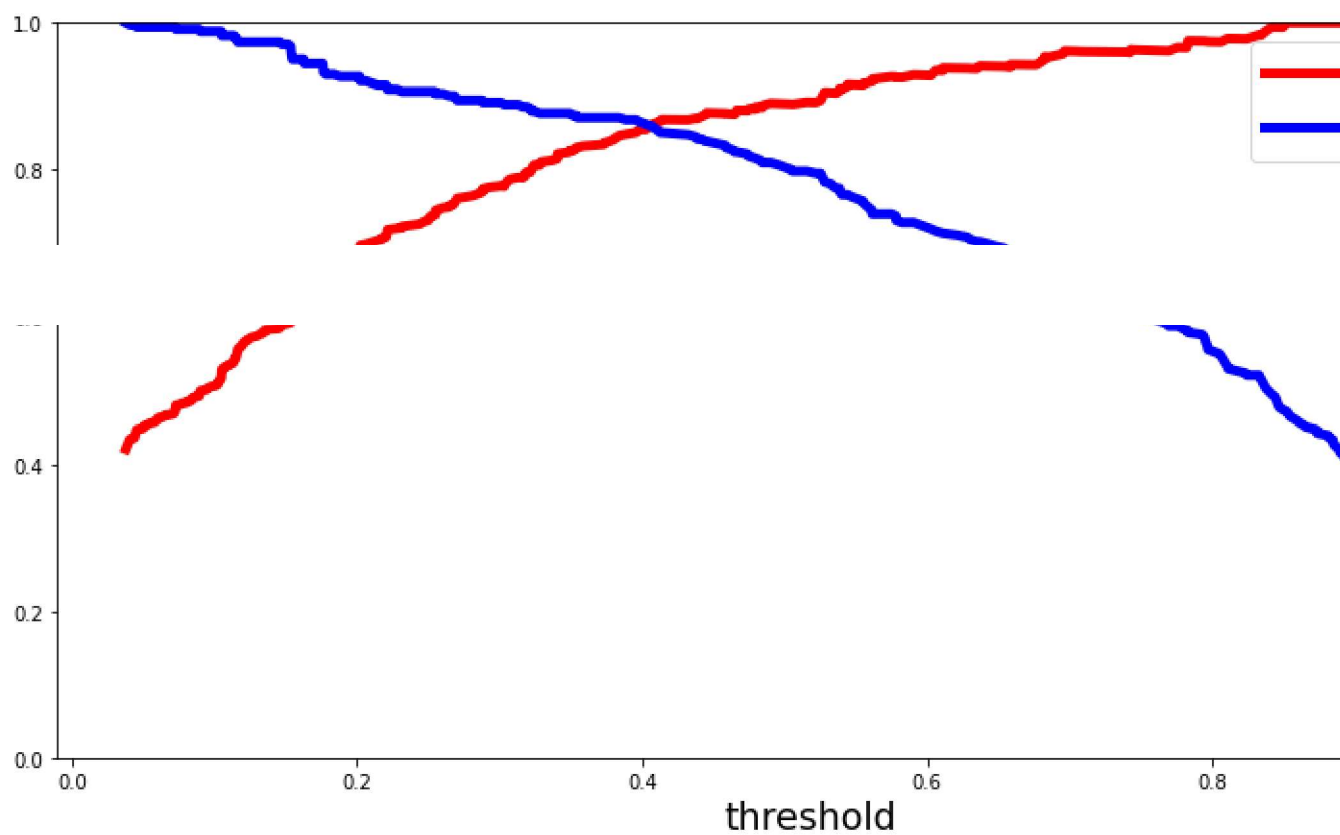
```
    plt.ylim([0, 1])
```

```
plt.figure(figsize=(14, 7))
```

```
plot_precision_and_recall(precision, recall, threshold)
```

```
plt.show()
```





[Colab paid products](#) - [Cancel contracts here](#)

✓ 0s completed at 7:37 PM

● ✕